



CS60010: Deep Learning

Spring 2023


Sudeshna Sarkar

Multilayer Perceptron - Introduction

Sudeshna Sarkar

20 Jan 2023

Linear Models




$$\mathbf{w} = [\theta_1 \ \theta_2 \ \dots \ \theta_d]^T \text{ and } \mathbf{x} = [x_1 \ x_2 \ \dots \ x_d]^T$$

$$\mathbf{z} = b + \sum_{i=1}^d \theta_i x_i = [\boldsymbol{\theta}^T b] \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}$$

$$\mathbf{y} = g(\mathbf{z})$$

- Regression: $g(\mathbf{z}) = \mathbf{z}$
- Classification:
 - Binary: $g(\mathbf{z}) = \sigma(\mathbf{z}) = \frac{1}{1+e^{-\mathbf{z}}}$
 - Multi-class

multilayer perceptrons




Intuition:-

Hidden Layer: Extracts better representation of the input data

Output layer: Does the classification

Neuron



$$\mathbf{w} = [w_1 \ w_2 \ \dots \ w_d]^T \text{ and } \mathbf{x} = [x_1 \ x_2 \ \dots \ x_d]^T$$

$$\mathbf{z} = b + \sum_{i=1}^d w_i x_i = [\mathbf{w}^T \mathbf{b}] \begin{bmatrix} x \\ 1 \end{bmatrix}$$
$$\mathbf{a} = g(\mathbf{z})$$

Terminologies:-

\mathbf{x} : input, \mathbf{w} : weights, \mathbf{b} : bias

a : pre-activation (input activation)

g : activation function

y : activation (output activation)

Output Units: Linear

$$\hat{y} = w^T a + b$$

Used to produce the mean of a conditional Gaussian distribution:

$$p(\mathbf{y} | \mathbf{x}) = \mathcal{N}(\mathbf{y}; \hat{\mathbf{y}}, \sigma)$$

Maximizing log-likelihood \Rightarrow Minimizing squared error

Output Units: Sigmoid

$$\hat{y} = \sigma(w^T a + b)$$

$$\begin{aligned} J(\theta) &= -\log p(y|x) \\ &= -\log \sigma((2y - 1)(w^T a + b)) \end{aligned}$$

Output Softmax Units



Need to produce a vector $\hat{\mathbf{y}}$ with $\hat{y}_i = p(y = i|x)$

$$\text{softmax}(z)_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

$$\log \text{softmax}(z)_i = z_i - \log \sum_j \exp(z_j)$$

Loss Functions




Regression

- Squared error: $L(y, \hat{y}) = (y - \hat{y})^2$

Classification

- Cross entropy: $L(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_k y_k \log \hat{y}_k$

Artificial Neuron – hidden unit



$$\mathbf{w} = [w_1 \ w_2 \ \dots \ w_d]^T \text{ and } \mathbf{x} = [x_1 \ x_2 \ \dots \ x_d]^T$$

$$\mathbf{z} = b + \sum_{i=1}^d w_i x_i = [\mathbf{w}^T \ b] \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}$$
$$\mathbf{a} = g(\mathbf{z})$$

Terminologies

\mathbf{x} : input, \mathbf{w} : weights, \mathbf{b} : bias

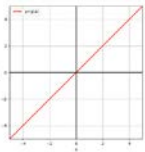
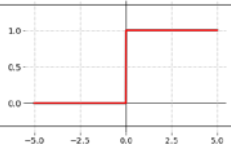
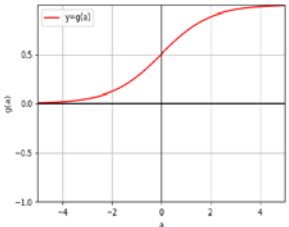
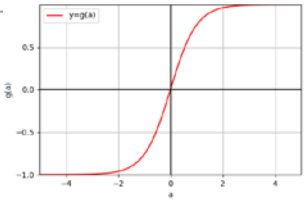
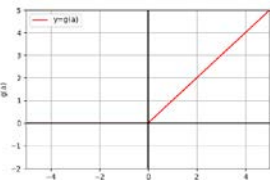
z : pre-activation (input activation)

g : activation function


\mathbf{a} : activation at hidden units

Common Activation Functions




Name	Function	Gradient	Graph
Linear	a	1	
Binary step	$\text{sign}(a)$	$g'(a) = \begin{cases} 0, & a \neq 0 \\ NA, & a = 0 \end{cases}$	
Sigmoid	$\sigma(a) = \frac{1}{1 + \exp(-a)}$	$g'(a) = g(a)(1 - g(a))$	
Tanh	$\tanh(a) = \frac{\exp(a) - \exp(-a)}{\exp(a) + \exp(-a)}$	$g'(a) = 1 - g^2(a)$	
ReLU	$g(a) = \max(0, a)$	$g'(a) = \begin{cases} 1, & a \geq 0 \\ 0, & a < 0 \end{cases}$	

Common Activation Functions




Linear activation function

- $g(a) = a$
- Unbounded
- $g'(a) = 1$




Sigmoid activation function

- $g(a) = \sigma(a) = \frac{1}{1+\exp(-a)}$
- Bounded (0, 1)
- Always positive
- $g'(a) = g(a)(1 - g(a))$



tanh activation function

- $g(a) = \tanh(a) = \frac{\exp(a) - \exp(-a)}{\exp(a) + \exp(-a)}$
- Bounded (-1, 1)
- Can be positive or negative
- $g'(a) = 1 - g^2(a)$



ReLU activation function

- $g(a) = \max(0, a)$
- Bounded below by 0
- But not upper-bounded
- $g'(a) = \begin{cases} 1, & a \geq 0 \\ 0, & a < 0 \end{cases}$

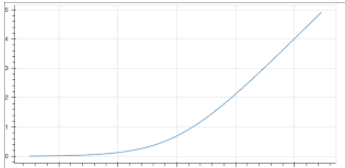
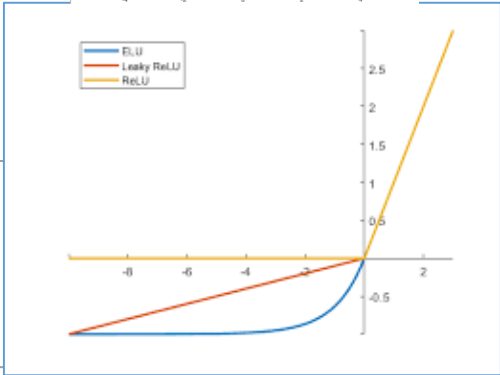
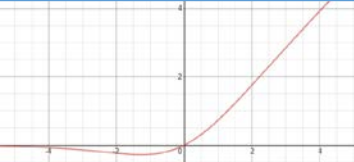
Activation Functions for Hidden Nodes





Name	Function	Gradient	Graph
Sigmoid	$\sigma(z) = \frac{1}{1 + \exp(-z)}$	$g'(z) = g(z)(1 - g(z))$	
Tanh	$\tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$	$g'(z) = 1 - g^2(z)$	
ReLU	$g(z) = \max(0, z)$	$g'(z) = \begin{cases} 1, & z \geq 0 \\ 0, & z < 0 \end{cases}$	
Softplus	$g(z) = \ln(1 + e^z)$		

More activation functions



Name	Function	Gradient	Graph
Softplus	$g(z) = \ln(1 + e^z)$	$g'(z) = \frac{1}{1 + e^{-z}}$	
Leaky Relu	$g(z) = \begin{cases} \alpha z, & z < 0 \\ z, & z \geq 0 \end{cases}$	$g'(z) = \begin{cases} \alpha, & z < 0 \\ 1, & z \geq 0 \end{cases}$	
ELU	$g(z) = \begin{cases} z, & z > 0 \\ \alpha(e^z - 1), & z \leq 0 \end{cases}$	$g'(z) = \begin{cases} 1, & z > 0 \\ \alpha(e^z), & z \leq 0 \end{cases}$	
swish	$g(z) = z \cdot \sigma(\beta z)$	$g'(z) = \beta g(\beta z) + \sigma(\beta z)(1 - \beta g(\beta z))$	

Rectified Linear Units



- Activation function: $g(z) = \max\{0, z\}$ with $z \in \mathbb{R}$
- Give large and *consistent* gradients when active
- **Good practice:** Initialize **b** to a small positive value (e.g. 0.1) Ensures units are initially active for most inputs and derivatives can pass through

Positives:

- Gives large and *consistent* gradients (does not saturate) when active
- Efficient to optimize, converges much faster than sigmoid or tanh

Negatives:

- Non zero centered output
- Units "die" i.e. when inactive they will never update

Swish



$$f(x) = x * \sigma(\beta x)$$




Figure 4: The Swish activation function.




Figure 5: First derivatives of Swish.


GELU



$$\text{gelu}(x) = x \Pr(X \leq x) \quad X \sim \mathcal{N}(0,1)$$

$$\text{gelu}(x) = \frac{1}{2} x \left(1 + \text{erf} \left(\frac{x}{\sqrt{2}} \right) \right)$$

GELU function and it's Derivative





Feedforward Networks and Backpropagation


Introduction




- **Goal:** Approximate some unknown ideal function $f^*: X \rightarrow Y$
- **Ideal classifier:** $y = f^*(x)$ for (x, y)
- **Feedforward Network:** Define parametric mapping $y = f(x; \theta)$
- **Learn** parameters θ to get a good approximation to f^* from training data
- Function f is a composition of many different functions e.g.

$$f(x) = f^3 \left(f^2 \left(f^1(x) \right) \right)$$

- **Training:** Optimize θ to drive $f(x; \theta)$ closer to $f^*(x)$
 - Only specifies the output of the *output layers*
 - Output of intermediate layers is not specified by D , hence the nomenclature *hidden layers*
- **Neural:** Choices of $f^{(i)}$'s and layered organization, loosely inspired by neuroscience



Universality and Depth



- First layer:

$$a^1 = g^1 \left(W^{1T} x + b^1 \right)$$
$$a^2 = g^2 \left(W^{2T} a^1 + b^2 \right)$$

- How do we decide depth, width?
- In theory how many layers suffice?

Universality



- Theoretical result [Cybenko, 1989]: 2-layer net with linear output with some squashing non-linearity in hidden units can approximate any continuous function over compact domain to arbitrary accuracy (given enough hidden units!)
- Implication: Regardless of function we are trying to learn, we know a large MLP can represent this function
- But not guaranteed that our training algorithm will be able to learn that function
- Gives no guidance on how large the network will be (exponential size in worst case)

A visual proof that neural nets can compute any function



- See Chapter 4 of **Neural Networks and Deep Learning** by Michael Nielsen
- <http://neuralnetworksanddeeplearning.com/chap4.html>

Advantages of Depth

