



CS60010: Deep Learning

Spring 2023

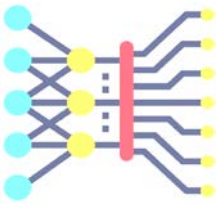
Sudeshna Sarkar

RNN Part 3 and Attention

Sudeshna Sarkar

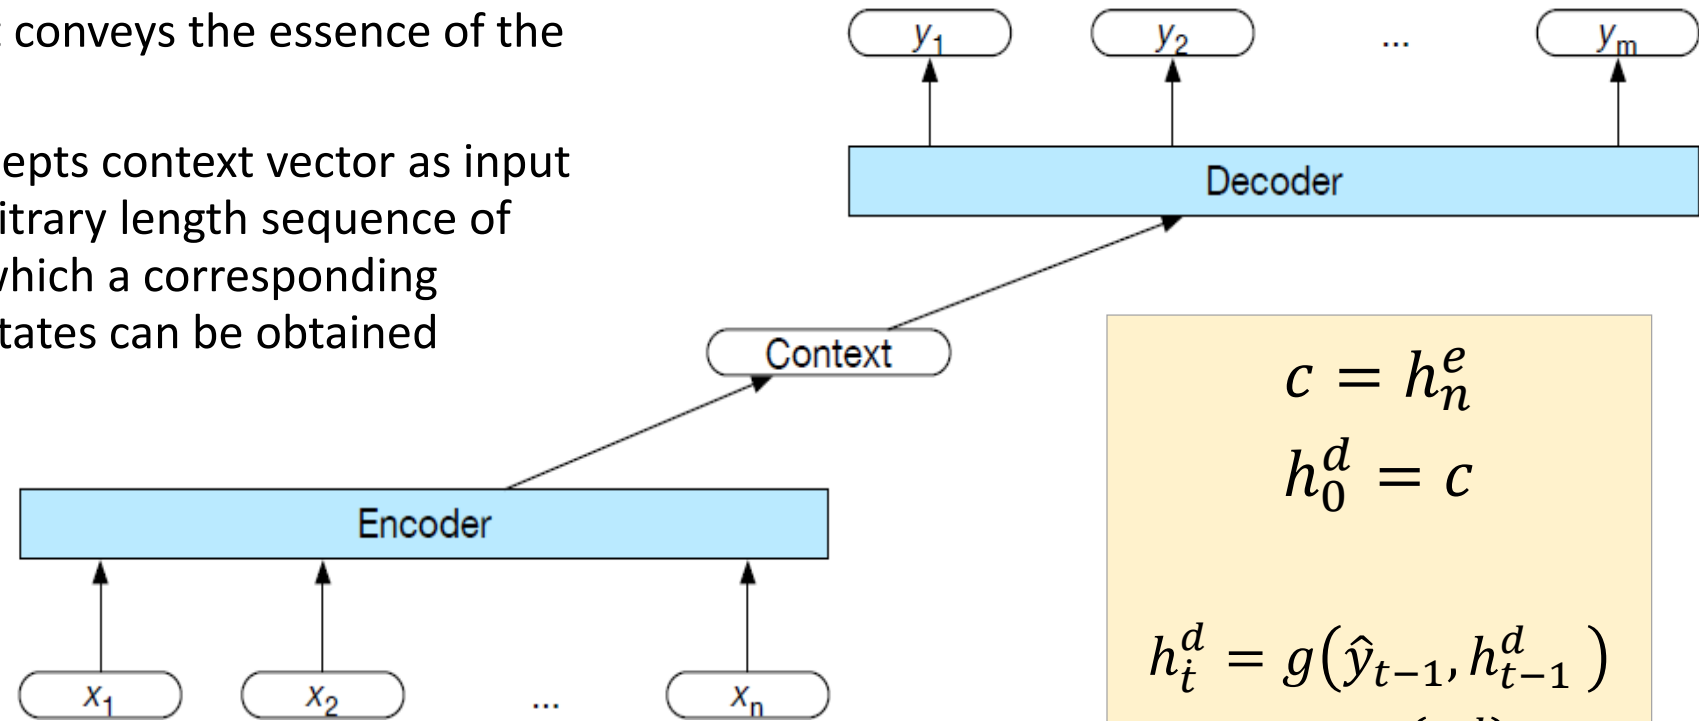
2 Mar 2023

Encoder-decoder networks



- An **encoder** that accepts an input sequence and generates a corresponding sequence of contextualized representations
- A **context vector** that conveys the essence of the input to the decoder
- A **decoder**, which accepts context vector as input and generates an arbitrary length sequence of hidden states, from which a corresponding sequence of output states can be obtained

simple RNNs, LSTMs, GRUs,
stacked Bi-LSTMs widely used
CNN



$$c = h_n^e$$

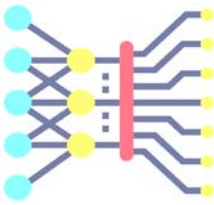
$$h_0^d = c$$

$$h_t^d = g(\hat{y}_{t-1}, h_{t-1}^d)$$

$$z_t = f(h_t^d)$$

$$y_t = \text{soft max}(z_t)$$

Decoder Weaknesses



The context vector c *only* available at the beginning of the generation process.

- Its influence became less-and-less important as the output sequence was generated.
- Solution: Make c available at each step in the decoding process,
 1. when generating the hidden states in the deocoder
 2. while producing the generated output.

$$c = h_n^e$$

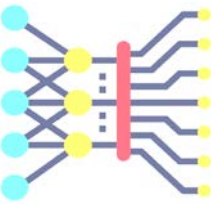
$$h_0^d = c$$

$$h_t^d = g(\hat{y}_{t-1}, h_{t-1}^d, c)$$

$$z_t = f(h_t^d)$$

$$y_t = \text{soft max}(\hat{y}_{t-1}, z_t, c)$$


Choosing the best output




- For neural generation, we can sample from the softmax distribution.
- In MT where we're looking for a specific output sequence, sampling isn't useful.
- Greedy Decoding: we choose the most likely output at each time step by taking the argmax over the softmax output

$$\hat{y} = \operatorname{argmax} P(y_i / y_{<i})$$


Greedy decoding



•




greedy decoding : take most probable word on each step

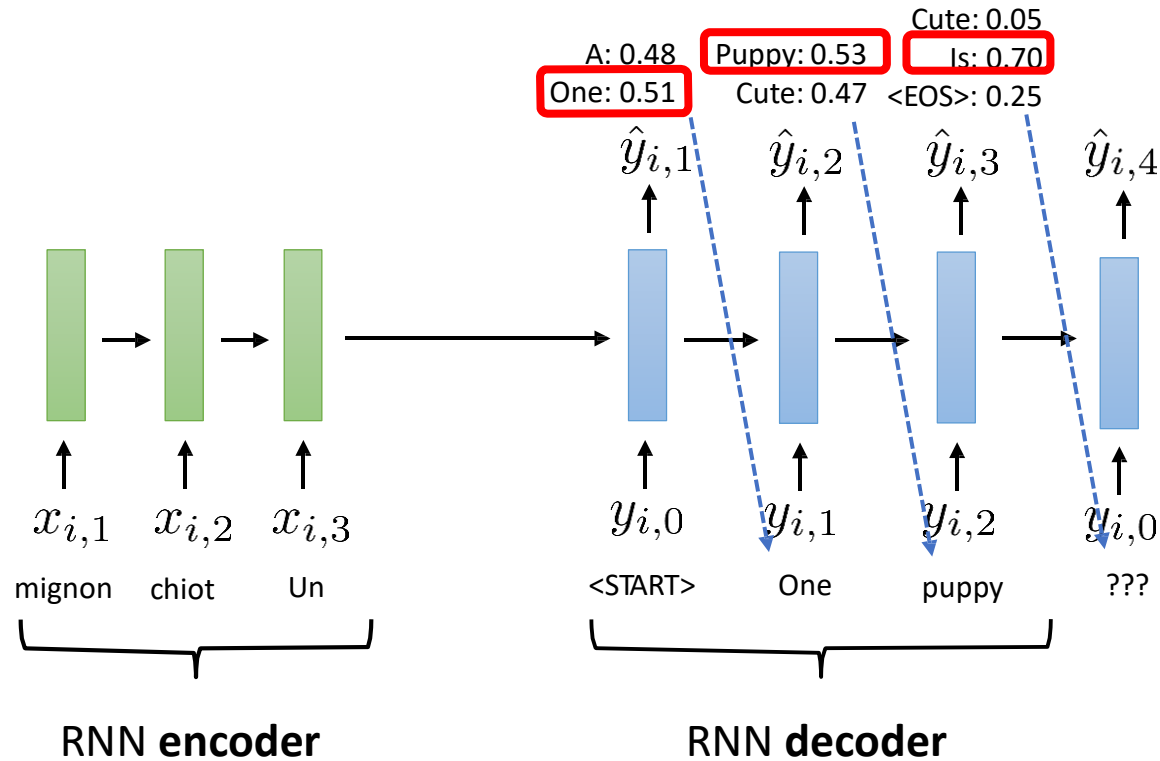


Decoding with beam search


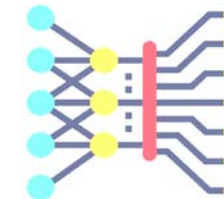
Decoding the most likely sequence



notice we feed
this in reverse

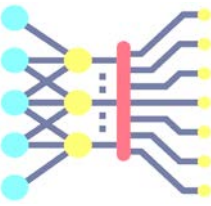


What we *should* have done




$$p(y_{i,1:T_y} | x_{i,1:T}) = \prod_{t=1}^{T_y} p(y_{i,t} | x_{i,1:T}, y_{i,0:t-1})$$

probabilities at each time step




How many possible decodings are there?



for M words, in general there are M^T sequences
of length T

any one of these might be the optimal one!


Decoding is a **search** problem



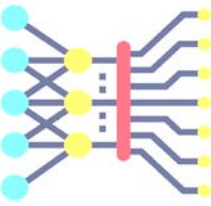
02-Mar-23

We could use *any* tree search
algorithm But exact search in this case
is **very** expensive
The **structure** of this problem makes
some simple **approximate search**
methods work **very well**

Beam search decoding



- Core idea: On each step of decoder, keep track of the ***k* most probable** partial translations (which we call ***hypotheses***)
 - *k* is the **beam size** (in practice around 5 to 10)
- A hypothesis y_1, \dots, y_t has a score which is its log probability:
$$\text{score}(y_1, \dots, y_t) = \log P_{\text{LM}}(y_1, \dots, y_t | x) = \sum_{i=1}^t \log P_{\text{LM}}(y_i | y_1, \dots, y_{i-1}, x)$$
 - Scores are negative, and higher score is better
 - We search for high-scoring hypotheses, tracking top *k* on each step
- Beam search is **not guaranteed** to find optimal solution
- But **much more efficient** than exhaustive search!



Decoding with approximate search

Basic intuition: while choosing the **highest-probability** word on the first step may not be optimal, choosing a **very low-probability** word is very unlikely to lead to a good result

Equivalently: we can't be greedy, but we can be *somewhat greedy*.


Beam search intuition: store the **k** best sequences **so far**, and update each of them.

special case of **k** = 1 is just

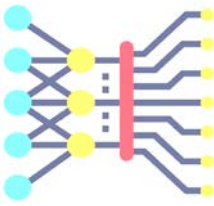
greedy decoding often use **k**

around 5-10

Decoding is a **search** problem



Beam search example



$$p(y_{i,1:T_y} | x_{i,1:T}) = \prod_{t=1}^{T_y} p(y_{i,t} | x_{i,1:T}, y_{i,0:t-1})$$

$$\log p(y_{i,1:T_y} | x_{i,1:T}) = \sum_{t=1}^{T_y} \log p(y_{i,t} | x_{i,1:T}, y_{i,0:t-1})$$

in practice, we **sum up** the log probabilities as we go (to avoid underflow)

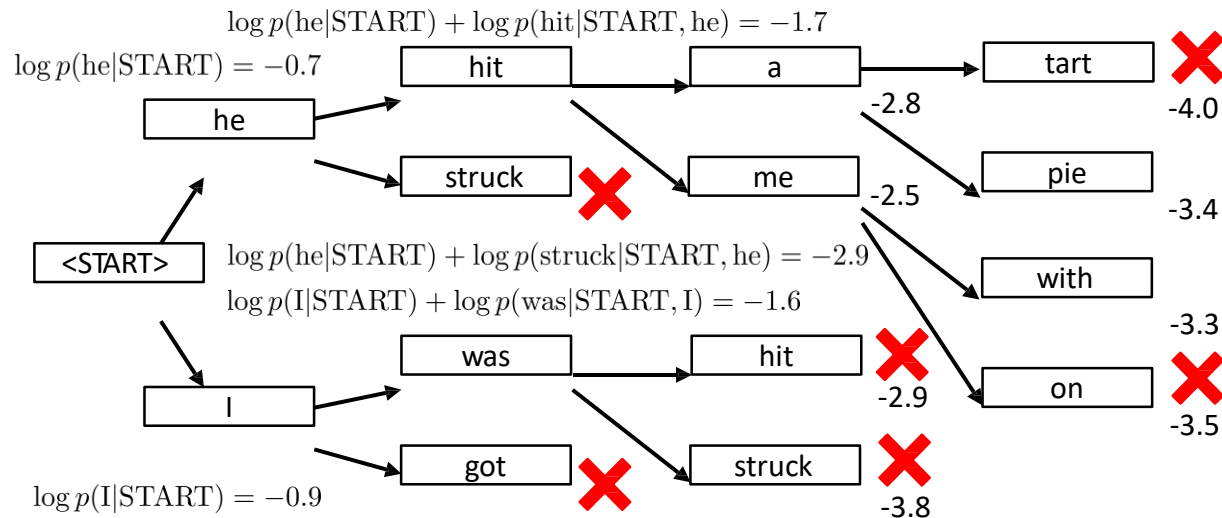
Example

k = 2 (track the 2 most likely hypotheses)

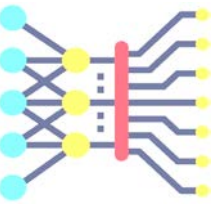
translate (Fr->En): il a m'entarté

(he hit me with a pie)

no perfectly equivalent English word, makes this hard



...and many other choices with lower log-prob



Beam search summary

$$\log p(y_{i,1:T_y} | x_{i,1:T}) = \sum_{t=1}^{T_y} \log p(y_{i,t} | x_{i,1:T}, y_{i,0:t-1})$$

there are k of these

at each time step t :

1. for each hypothesis $y_{1:t-1,i}$ that we are tracking:


find the top k tokens $y_{t,i,1}, \dots, y_{t,i,k}$

very easy, we get this from the softmax log-probs

2. sort the resulting k^2 length t sequences by their *total* log-probability

3. keep the top k

4. advance each hypothesis to time $t + 1$

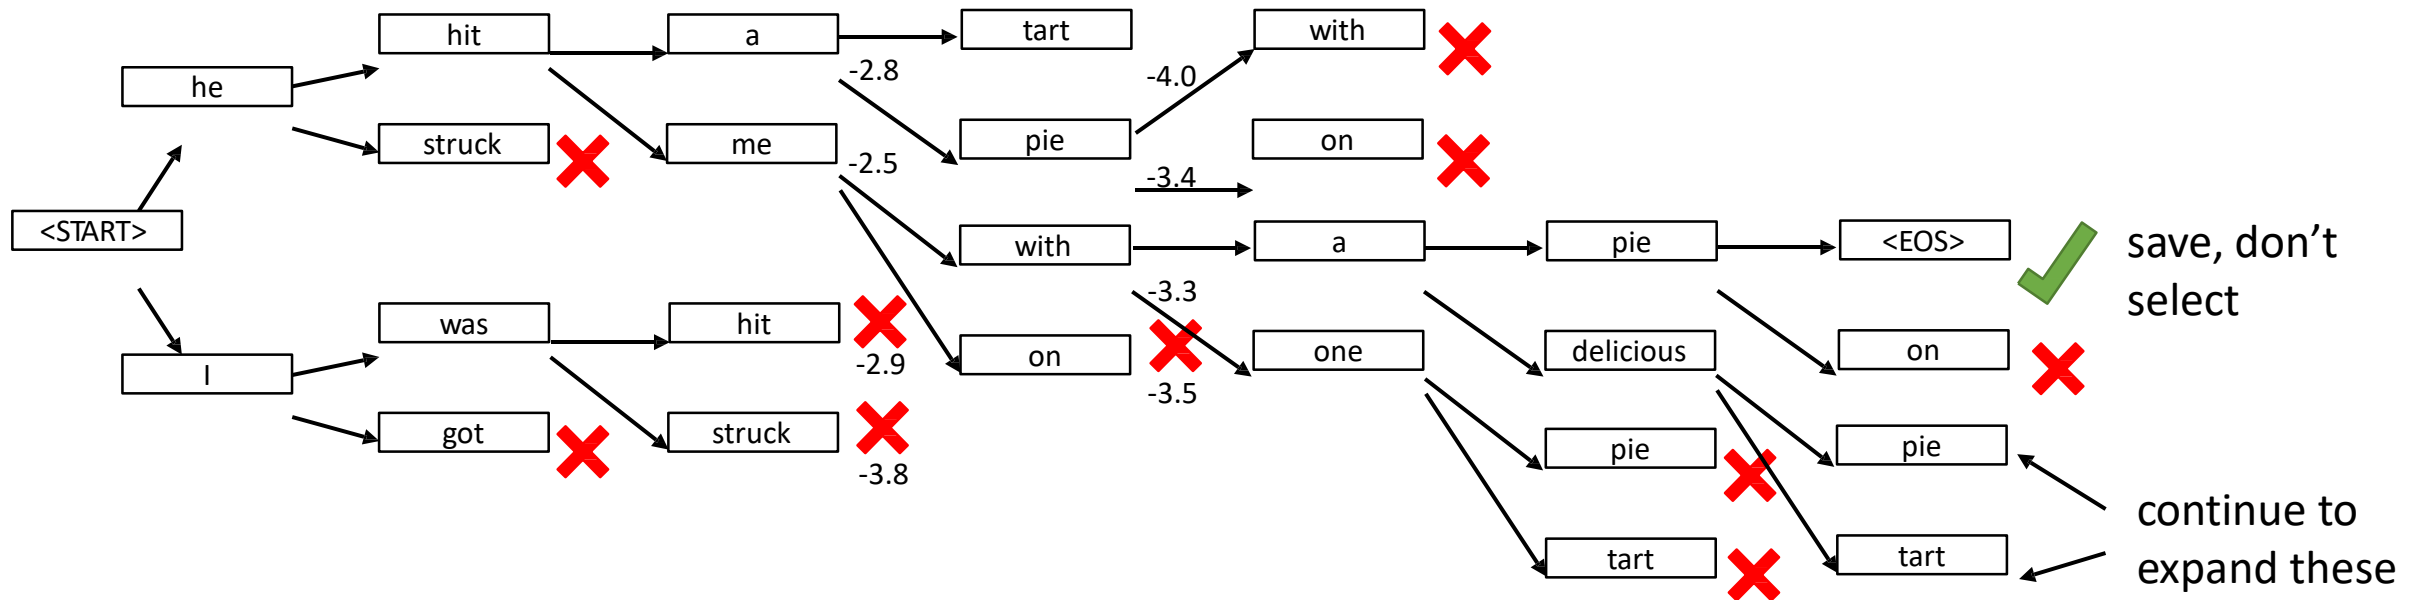


When do we stop decoding?

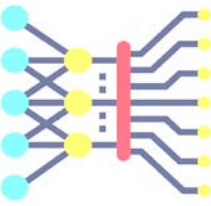
Suppose one of the highest-scoring hypotheses ends in <END>

Save it, along with its score, but do **not** pick it to expand further (there is nothing to expand)

Keep expanding the **k** remaining best hypotheses



Continue until either some cutoff length **T** or until we have **N** hypotheses that end in <EOS>



Which sequence do we pick?

At the end we might have something like this:

he hit me with a pie he $\log p = -4.5$

threw a pie $\log p = -3.2$

I was hit with a pie that he threw $\log p = -7.2$

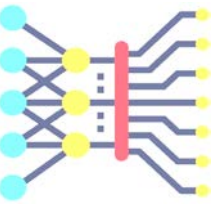
$$\log p(y_{i,1:T}|x_{i,1:T}) = \sum_{t=1}^T \log p(y_{i,t}|x_{i,1:T}, y_{i,0:t-1})$$

Problem: $p < 1$ **always**, hence $\log p < 0$ **always**

The **longer** the sequence the **lower** its total score (more negative numbers added together)

Simple “fix”: just divide by sequence length

$$\text{score}(y_{i,1:T}|x_{i,1:T}) = \frac{1}{T} \sum_{t=1}^T \log p(y_{i,t}|x_{i,1:T}, y_{i,0:t-1})$$




Beam search summary

$$\text{score}(y_{i,1:T}|x_{i,1:T}) = \frac{1}{T} \sum_{t=1}^T \log p(y_{i,t}|x_{i,1:T}, y_{i,0:t-1})$$

at each time step t :

1. for each hypothesis $y_{1:t-1,i}$ that we are tracking:
find the top k tokens $y_{t,i,1}, \dots, y_{t,i,k}$
2. sort the resulting k^2 length t sequences by their *total* log-probability
3. save any sequences that end in EOS
4. keep the top k
5. advance each hypothesis to time $t + 1$ if $t < H$

return saved sequence with highest score



CS60010: Deep Learning

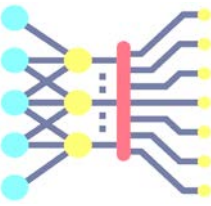
Spring 2023

Sudeshna Sarkar

Attention


Sudeshna Sarkar

3 Mar 2023



Encoder Decoder Model


Conditional Language Model



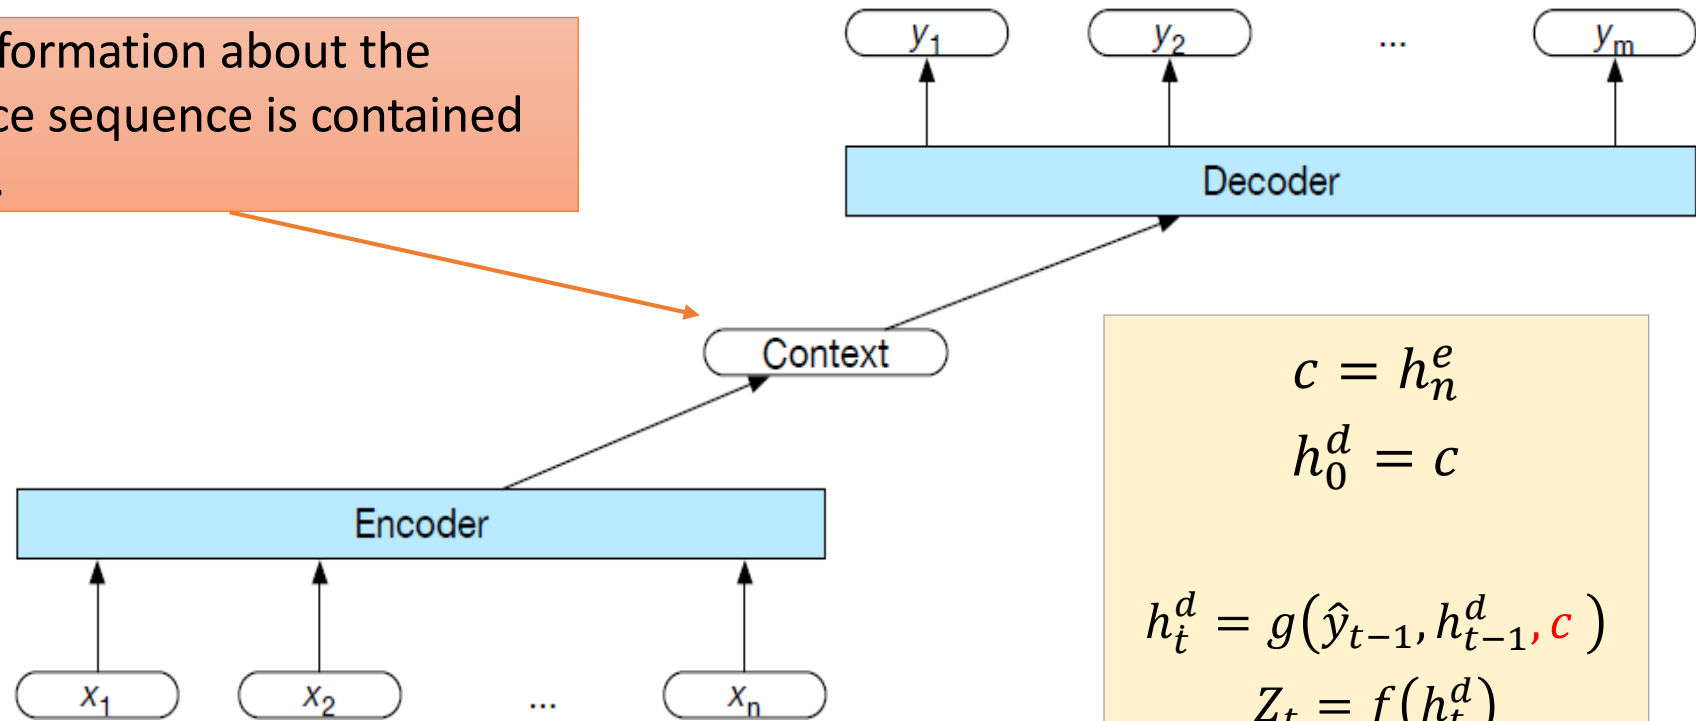
typically two **separate** RNNs (with different weights)

trained **end-to-end** on paired data (e.g., pairs of French & English sentences)

Encoder-decoder networks



all information about the source sequence is contained here.



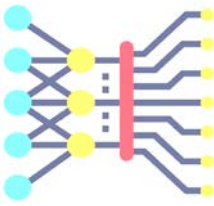
$$c = h_n^e$$
$$h_0^d = c$$

$$h_t^d = g(\hat{y}_{t-1}, h_{t-1}^d, c)$$

$$z_t = f(h_t^d)$$

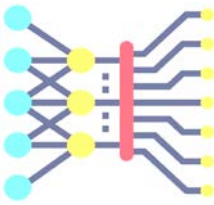
$$y_t = \text{softmax}(z_t)$$

Applications of Encoder- Decoder Networks



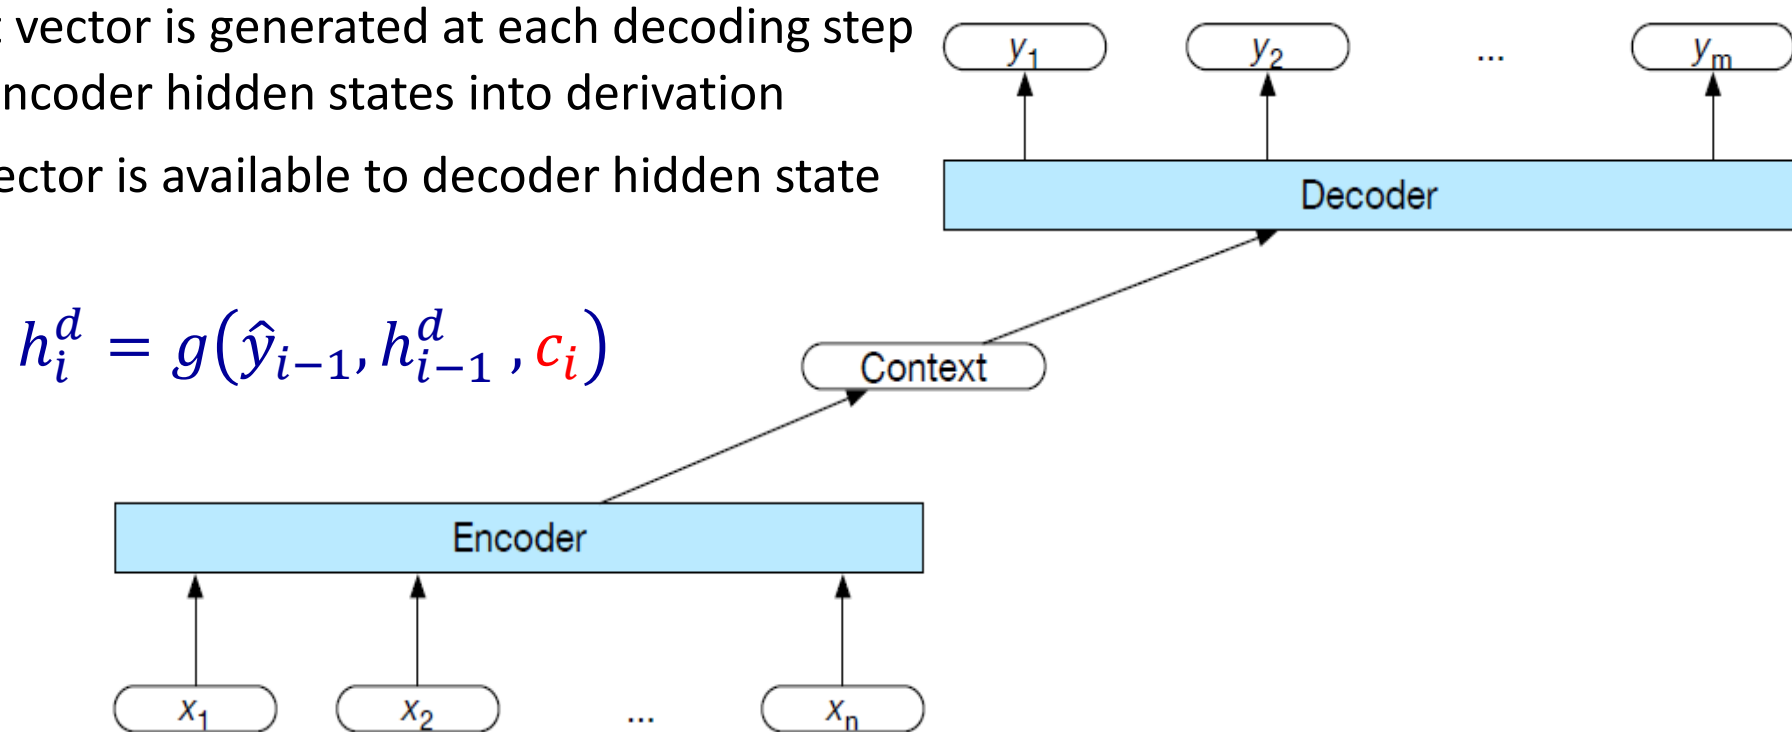
- Text summarization
- Text simplification
- Question answering
- Image captioning
- ...

Attention




- Replace the static context vector with one that is dynamically derived from the encoder hidden states at each point during decoding
- A new context vector is generated at each decoding step and takes all encoder hidden states into derivation
- This context vector is available to decoder hidden state calculations


$$h_t^d = g(\hat{y}_{t-1}, h_{t-1}^d, c)$$




The bottleneck problem



this forms a bottleneck



Idea: what if we could somehow “peek” at the source sentence while decoding?



To calculate c_i , first find relevance of each encoder hidden state to the decoder state $score(h_{i-1}^d, h_j^e)$ for each encoder state j

1. The *score* can simply be dot product

$$score(h_{i-1}^d, h_j^e) = h_{i-1}^d \cdot h_j^e$$

2. The score can also be parameterized with weights

$$score(h_{i-1}^d, h_j^e) = h_{i-1}^d W_s h_j^e$$



- Normalize with a softmax to create a vector of weights $\alpha_{i,j}$ that tells us the proportional relevance of each encoder hidden state j to the current decoder state i

$$\alpha_{i,j} = softmax(score(h_{i-1}^d, h_j^e) \forall j \in e)$$

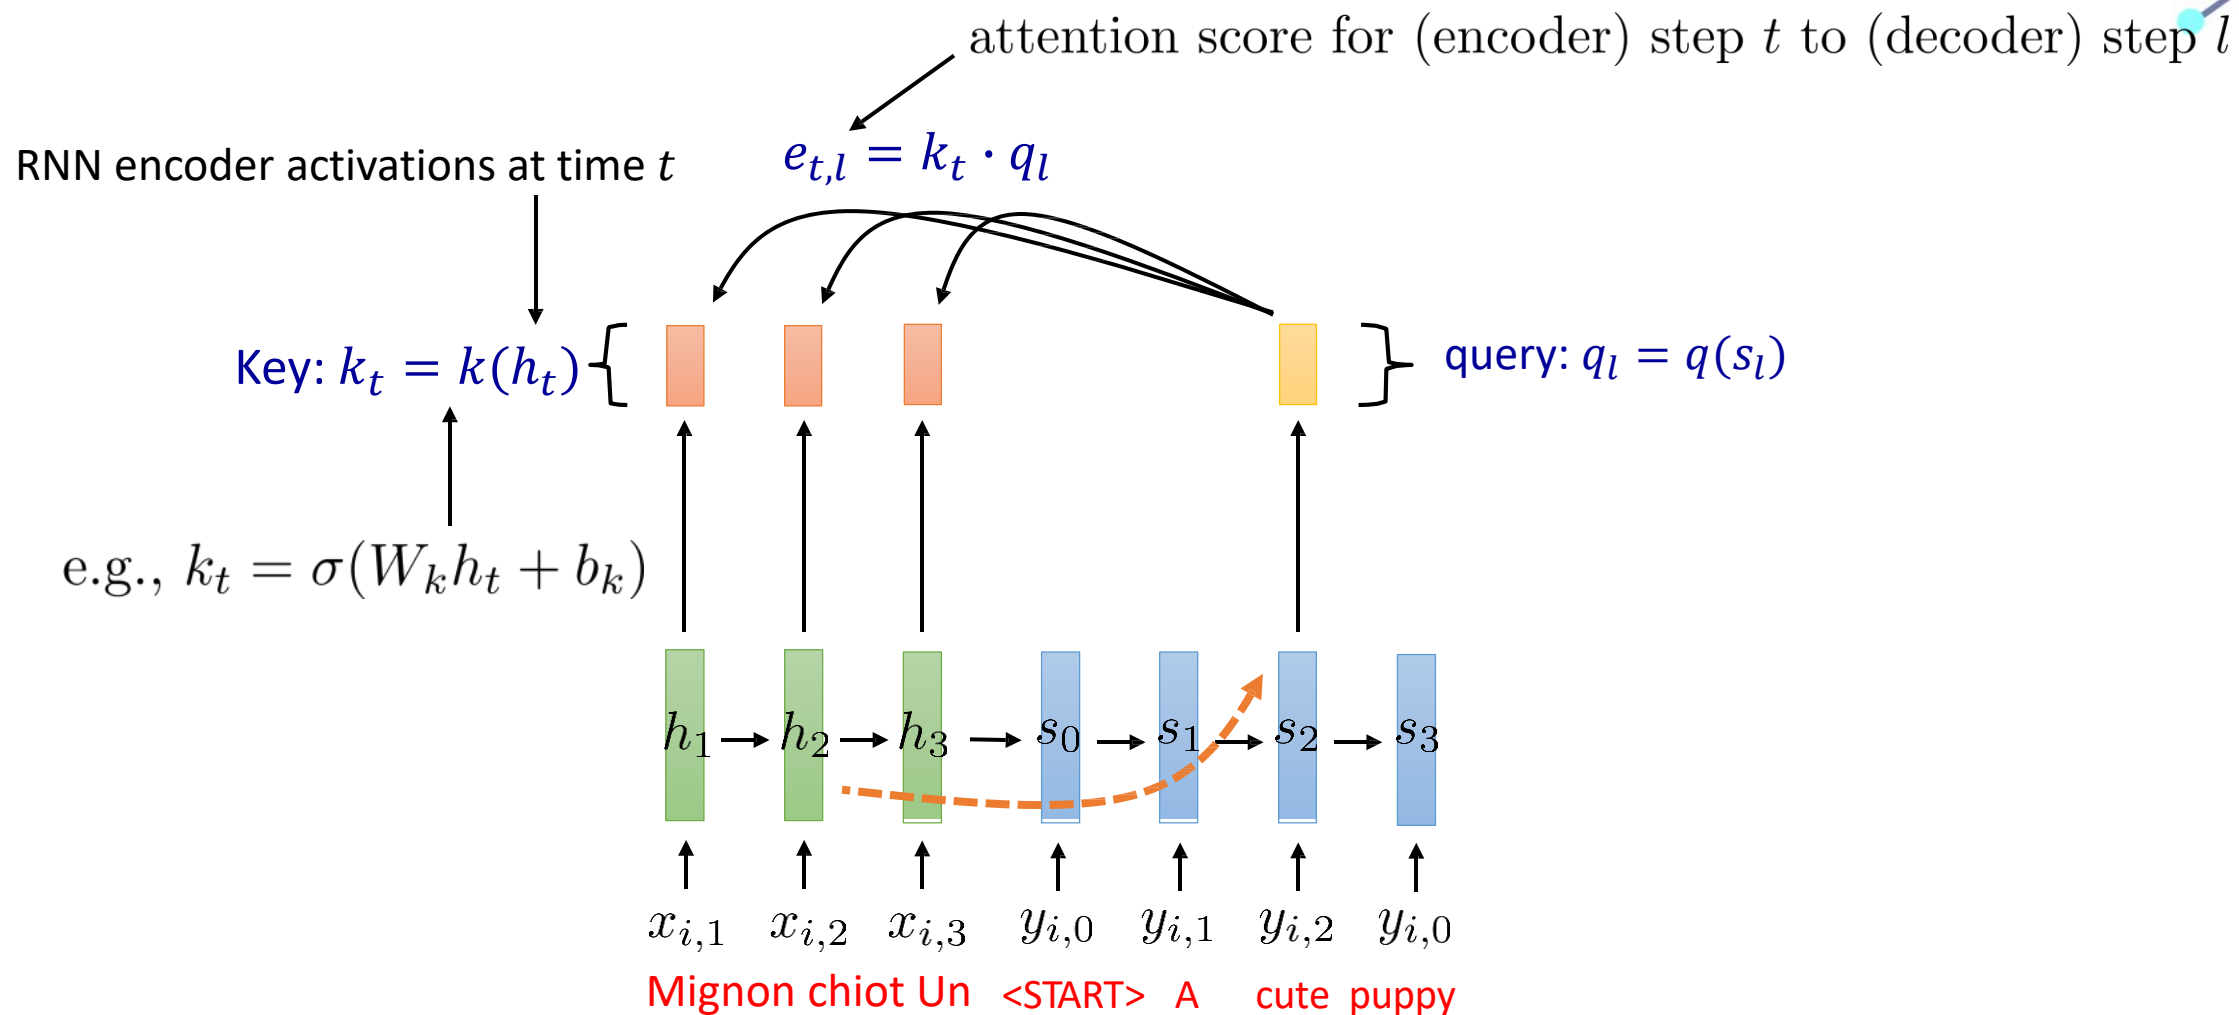

- Finally, context vector is the weighted average of encoder hidden states

$$c_i = \sum_j \alpha_{i,j} h_j^e$$



Can we “peek” at the input?




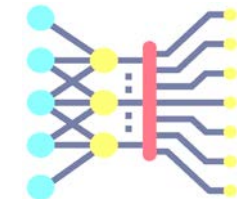
Attention



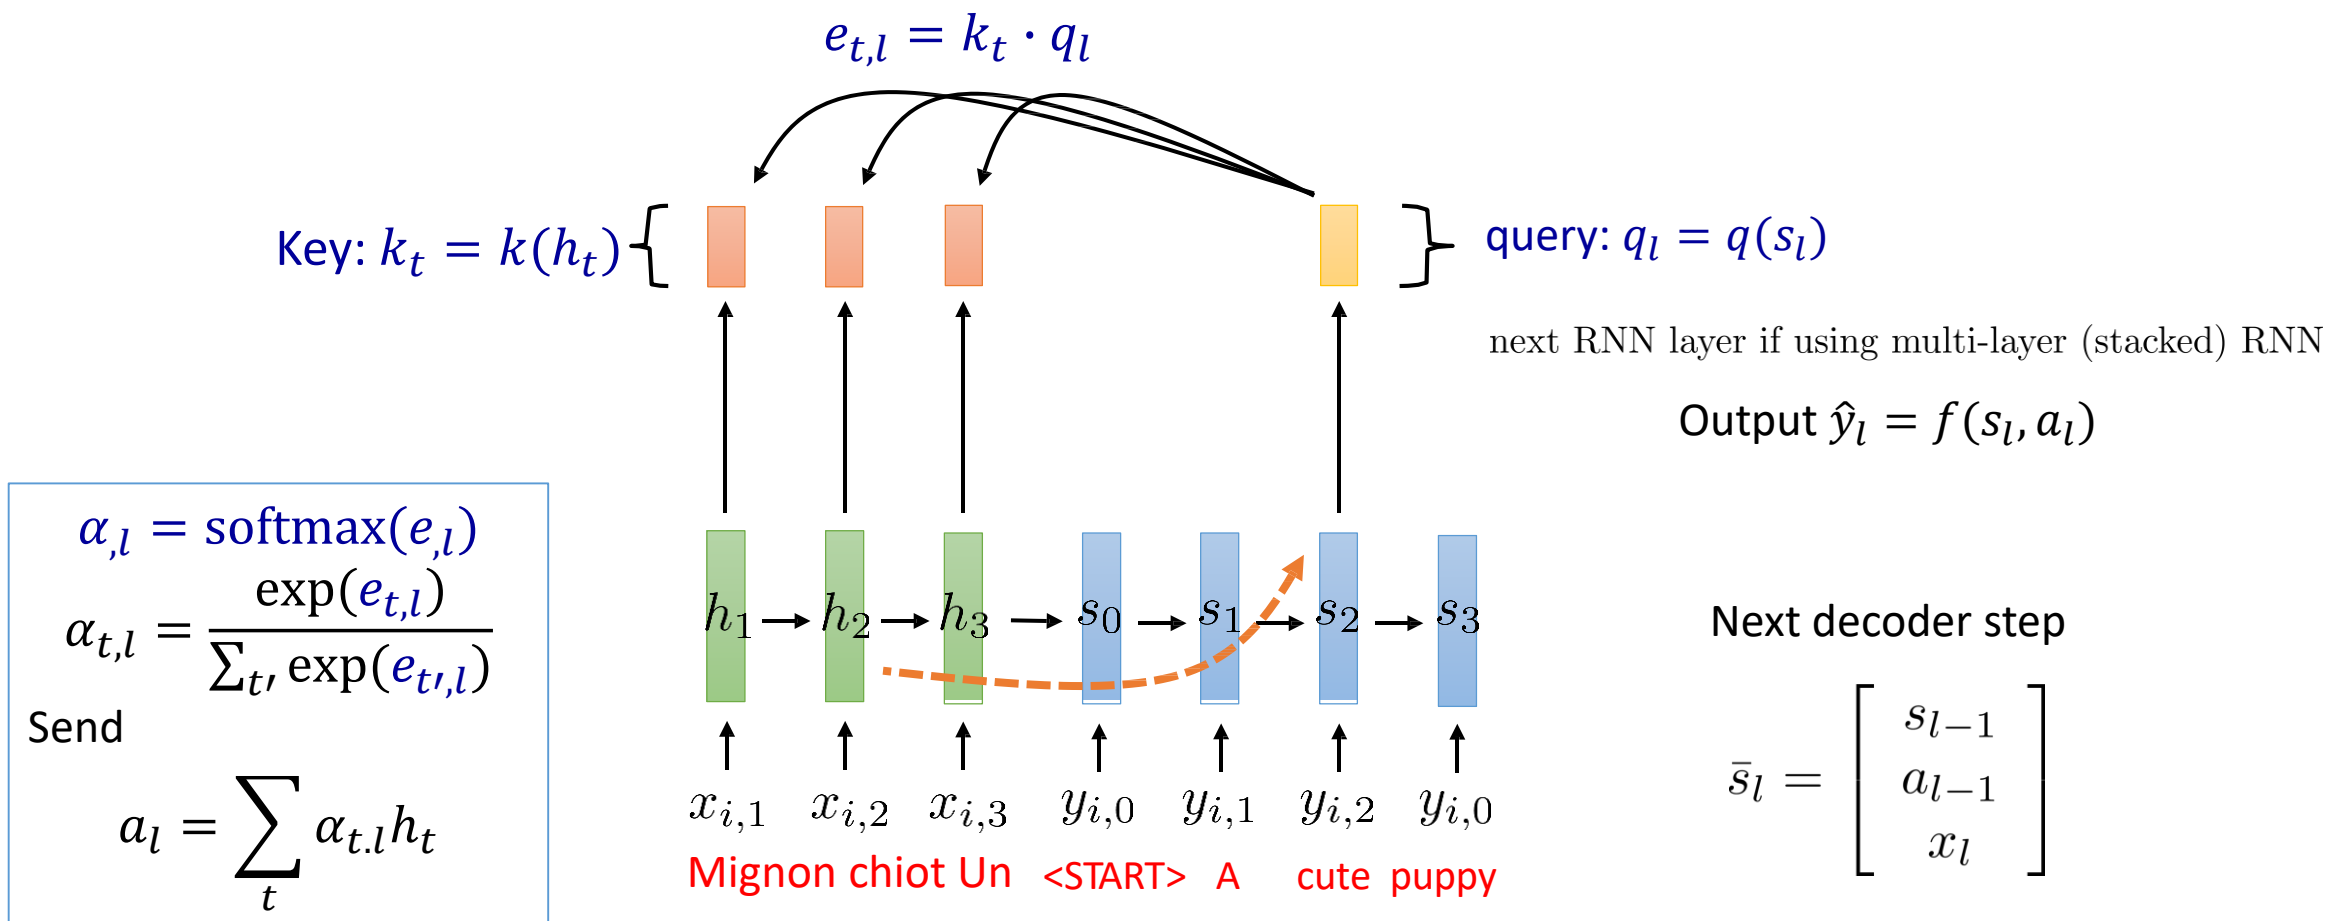

Attention





Attention



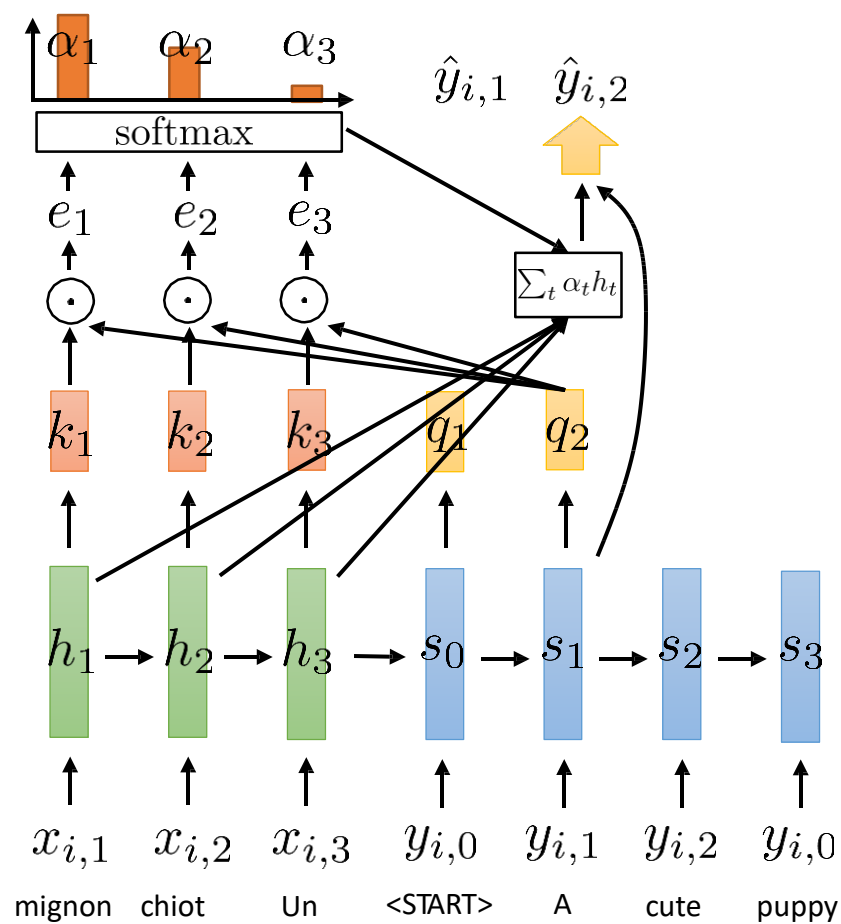

Attention





Attention Walkthrough (Example)



Attention Walkthrough (Example)



Attention Walkthrough (Example)



Attention Equations

Encoder-side:

$$k_t = k(h_t)$$

Decoder-side:

$$q_l = q(s_l)$$

$$e_{t,l} = k_t \cdot q_l$$

$$\alpha_{t,l} = \frac{\exp(e_{t,l})}{\sum_{t'} \exp(e_{t',l})}$$

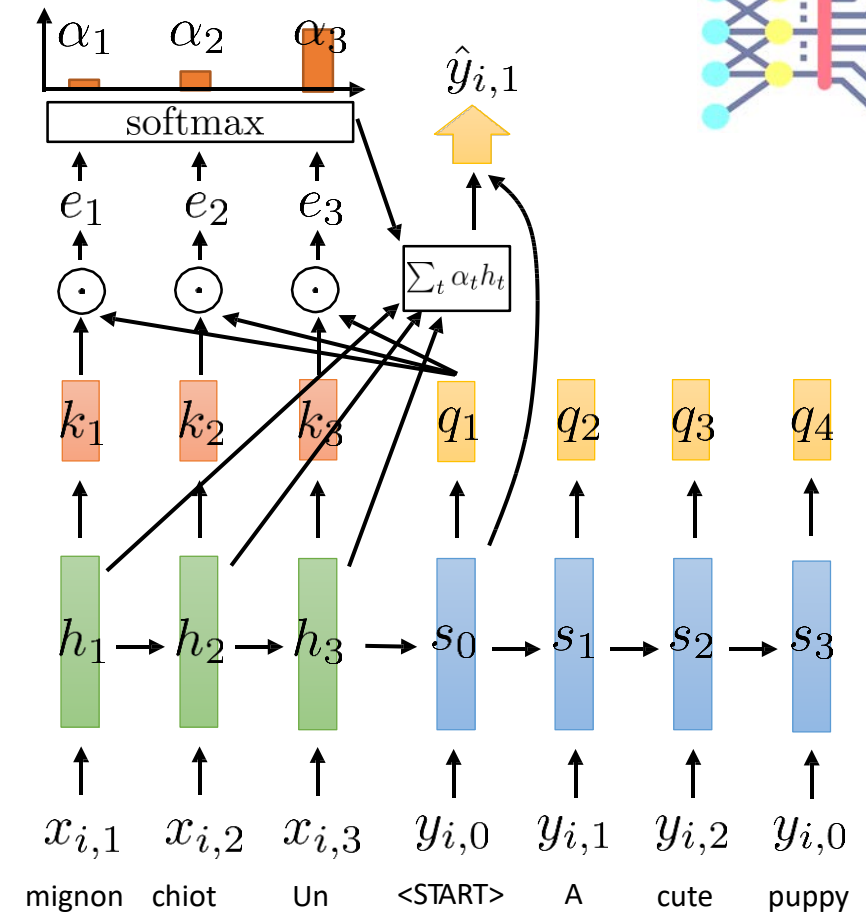
$$a_l = \sum_t \alpha_t h_t$$

Could use this in various ways:

concatenate to hidden state: $\begin{bmatrix} s_{l-1} \\ a_{l-1} \\ x_l \end{bmatrix}$

use for readout, e.g.: $\hat{y}_l = f(s_l, a_l)$

concatenate as input to next RNN layer



Attention Variants

Simple key-query choice: k and q are identity functions

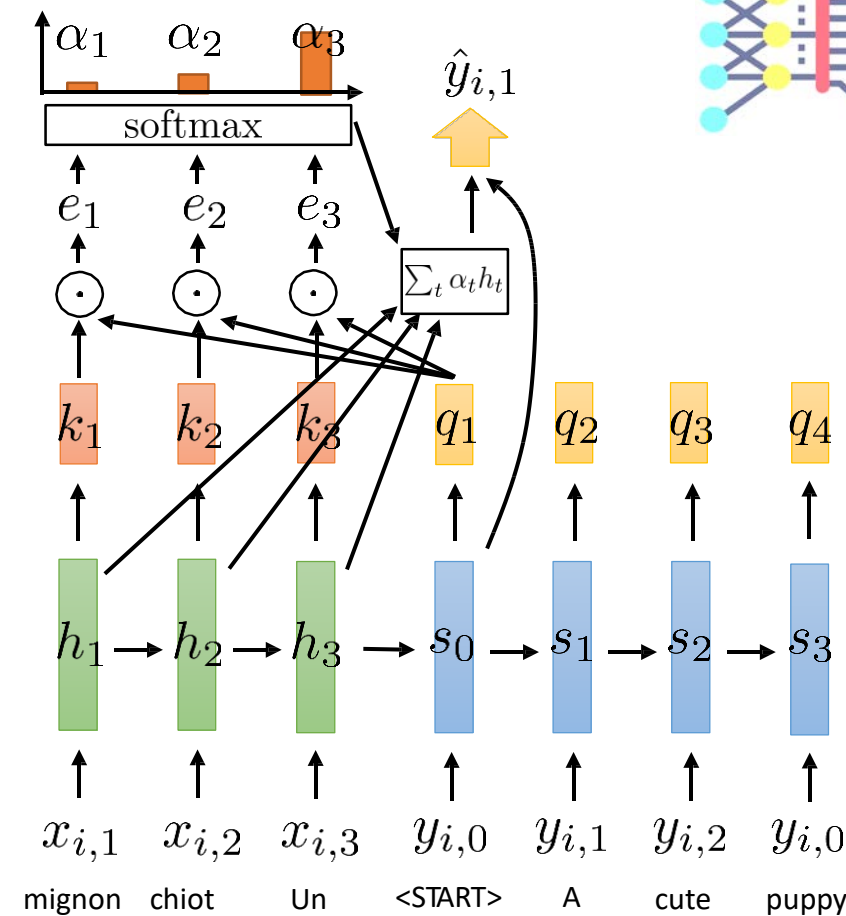
$$k_t = h_t \quad q_l = s_l$$

Decoder-side:

$$e_{t,l} = h_t \cdot s_l$$

$$\alpha_{t,l} = \frac{\exp(e_{t,l})}{\sum_{t'} \exp(e_{t',l})}$$

$$a_l = \sum_t \alpha_t h_t$$



Attention Variants

Linear multiplicative attention:


$$k_t = W_k h_t \quad q_l = W_q s_l$$

Decoder-side: just learn this matrix

$$e_{t,l} = h_t^T W_k^T W_q s_l = h_t^T W_e s_l$$

$$\alpha_{t,l} = \frac{\exp(e_{t,l})}{\sum_{t'} \exp(e_{t',l})}$$

$$a_l = \sum_t \alpha_t h_t$$



Attention Variants

Learned value encoding:

Encoder-side:

$$k_t = k(h_t)$$

Decoder-side:


$$q_l = q(s_l)$$

$$e_{t,l} = k_t \cdot q_l$$

$$\alpha_{t,l} = \frac{\exp(e_{t,l})}{\sum_{t'} \exp(e_{t',l})}$$

$$a_l = \sum_t \alpha_t v(h_t)$$

some learned function



Attention Summary

Every encoder step t produces a key k_t

Every decoder step l produces a query q_l

Decoder gets “sent” encoder activation h_t
corresponding to largest value of $k_t \cdot q_l$

actually gets $\sum_t \alpha_t h_t$

- Attention is **very** powerful, because now all decoder steps are connected to **all** encoder steps!
- Gradients are much better behaved ($O(1)$ propagation length)
- Becomes very important for very long sequences
- Bottleneck is much less important

